

Supplementary Information

An upstream placebo shows that reservoir siting explains the apparent drought buffering of seasonal reservoirs in Chile’s megadrought

Francisco Zambrano Bigiarini

2026-08-07


This Supplementary Information collects the supporting design, robustness, and diagnostic results referenced in the main text. Supplementary Figures S1–S3 document the drought forcing, the informativeness of the null, and the siting-confound decomposition; Supplementary Figures S4–S5 give the full study-area map and the matching diagnostics; Supplementary Figure S6 is the water-rights induced-demand test and Supplementary Figure S7 the barren-land ET-confound demonstration and Supplementary Figure S8 the non-parametric aridity common-support check; Supplementary Tables S1–S6 report the equivalence bounds, placebo comparability, positive control, decomposition ladder, storage-band trends, and the aridity-overlap sensitivity; Supplementary Tables S7–S9 add the fully non-parametric overlap re-match, the winsorization-threshold sensitivity, and the spatial-autocorrelation assessment with spatially-restricted inference; Supplementary Table S10 bounds the groundwater-substitution confound with the DGA well hydrograph network; Supplementary Table S11 tests reservoir-use heterogeneity (irrigation-only subset); Supplementary Table S12 tests the snowmelt confound on the differential estimand of the upstream/downstream placebo with standardized ERA5-Land snow-water-equivalent (snow-adjusted refits, a low-snow subsample, snow-year modulation, and an early/late stability split); Supplementary Table S13 shows the matching is insensitive to recomputing baseline aridity on the strictly pre-drought 1991–2009 window; Supplementary Table S14 reports the minimum detectable effect for the cropland pre-trend; Supplementary Table S15 collects the inter-basin spillover (cuenca disjointness) and administrative-demand checks, including the treated-versus-control closure test; Supplementary Table S16 translates the equivalence margin into physical streamflow volumes; Supplementary Table S17 reports the strict hard-balanced aridity sensitivity (effective control sample 19); and Supplementary Table S18 gives the rank-based tests on raw, uncapped water-rights volumes; Supplementary Table S19 the volumetric-scale (log-flow) streamflow checks; Supplementary Table S20 the metric and functional-form checks; and Supplementary Table S21 the relegated evapotranspiration outcomes; Supplementary Table S22 the geomorphological (alluvial) confound check; Supplementary Table S23 the carryover-capacity context for the treated reservoirs; and Supplementary Table S24 the early/late megadrought phase split of the streamflow buffering and placebo tests. Supplementary Table S25 is the direct within-basin test of the upstream/downstream placebo contrast; Supplementary Table S26 the permutation inference for the within-region (siting-ladder rung 4) residual; Supplementary Table S27 the permutation inference for the megadrought phase shift; Supplementary Tables S28 and S29 the non-linear (quadratic) and dry-tail transmission checks; Supplementary Table S30 the catchment-area comparability check; Supplementary Table S31 the elevation-adjusted placebo; Supplementary Table S32 the gauge-classification robustness of the contrast; Supplementary Table S33 the cumulative-scale (discharge-adjusted) placebo; Supplementary Table S34 the paired-placebo

contrast power; and Supplementary Table S35 the storage-capacity (carryover) heterogeneity of the buffering test. The downstream extraction-masking bound (the downstream-versus-upstream cropland gradient, the control transmission-slope sensitivity to local extraction, and the predicted masking magnitude) is Supplementary Table S36, the placebo net of a SPEI-by-local-cropland interaction is Supplementary Table S37, the sub-watershed SPEI exposure-misclassification bound is Supplementary Table S38, the strictly-paired carryover contrast is Supplementary Table S39, and the carryover-ratio denominator sensitivity (unregulated upstream flow) is Supplementary Table S40. Supplementary Table S41 decomposes the storage-band decline into a 2010 level step versus a continuing drift, and Supplementary Table S42 reports the carryover-claim hardening checks (the irrigation-only, hydropower-excluded contrast and the leave-one-out permutation inference). Supplementary Table S43 screens the upstream placebo for cascade-dam contamination against the 1,370-dam national inventory and re-runs the placebo without flagged gauges; Supplementary Table S44 re-fits the storage-band trends weighted by reservoir capacity and in absolute fleet volumes; and Supplementary Table S45 reports the hard-balance weight diagnostics and the naive and re-balanced leave-one-out sensitivities. Supplementary Table S46 validates the upstream/downstream gauge classification against the HydroRIVERS routed network and re-runs the placebo on network-connected gauges only; Supplementary Table S47 replaces the year fixed effects with a cubic national time trend so the temporal drought signal is retained; Supplementary Table S48 re-tests the aggregate and carryover contrasts under an elevation-partitioned permutation null; Supplementary Table S49 re-derives the streamflow results with SSI standardized on the 1990–2020 baseline; Supplementary Table S50 quantifies the drought-conditional structure of streamflow missingness; Supplementary Table S51 re-fits the storage-band trends on the coverage-stable reservoir subset; Supplementary Table S52 adds quadratic aridity-by-forcing interactions to the buffering model; and Supplementary Table S53 screens the matched controls against the 1,370-dam national inventory and re-fits the streamflow results with a re-balanced dam-free control pool. Supplementary Table S54 confirms every storage-band trend, step, and drift coefficient with a null-imposed wild cluster bootstrap, the small-cluster refinement appropriate to the 22 to 26 reservoir clusters. All items are regenerated by the `targets` pipeline.

Supplementary Figures

A sustained megadrought forcing, shared by dammed and control basins

Both track the same shock (dammed run slightly drier, a siting offset), so the estimand is the slope, not level




Supplementary Figure S1 | The drought forcing. (a) National mean SPEI-12 across the matched basins, 2000–2024; the negative (drought) excursion is filled and the post-2010 megadrought era shaded, with the 2019 and 2021 hyperdrought winters marked. The deficit is sustained and deep (below -1 for much of 2016–2024), giving the SPEI dose real dynamic range. (b) Entropy-balancing-weighted mean SPEI-12 for dammed versus matched-control basins: both co-move through the same megadrought (a common shock that year fixed effects absorb), with dammed basins running slightly drier, a siting offset that is precisely why the estimand is the differential deficit-to-impact slope rather than a level or calendar-time contrast.


The null is informative: large effects excluded, injected effects recovered

Rules out buffering beyond ~half the baseline slope; recovers a real one 1:1 (detected at $|\beta| \geq 0.30$)

a




b



Supplementary Figure S2 | The null is informative, not underpowered blindness. (a) Equivalence / minimum-detectable-effect bounds for the two streamflow outcomes: the observed slope gap (point) and its 90% (TOST) interval against the $\pm 25\%$ -of-baseline negligible region (grey). The 90% intervals exceed the negligible region, so formal equivalence to zero is not claimed, but they exclude attenuations larger than the minimum detectable effect (46% of baseline transmission for the intent-to-treat design, 59% downstream). (b) Positive control: a known buffering slope injected into the treated downstream gauges is recovered one-to-one by the estimator (points on the dashed expected-recovery line), the no-injection case is correctly not flagged, and randomization inference detects the effect (filled points) once the injected slope reaches -0.30 . The estimator would therefore reveal a large operational buffering effect if one existed.

Apparent buffering is a between-region siting confound

Survives matching and fixed effects; collapses only within climate region




Supplementary Figure S3 | Siting-confound decomposition ladder. Dammed-versus-control transmission-slope gap (treat:SPEI; <0 = apparent buffering) estimated by increasingly stringent estimators, for the streamflow SSI slope (top) and the two difference-in-differences proxy outcomes; bars are 95% cluster-robust confidence intervals and the design and within-region rows carry the design's randomization-inference p. Rungs: (1) naive pooled (no matching, no fixed effects, what a conventional dammed-versus-undammed study reports); (2) unit and time fixed effects, unweighted; (3) the design estimator (entropy-balance weights + fixed effects); (4) design with the baseline SPEI-to-outcome slope allowed to vary within climate region. For streamflow the apparent buffering survives rungs (1) to (3) near -0.18 (permutation-null throughout), then collapses toward zero at rung (4), and is equally strong at the unregulated upstream placebo (grey triangle): the signature of a between-region aridity confound, not regulation. The proxy outcomes carry no signal at any

rung.

A national matched design across Chile's aridity gradient

Dammed basins sit in the arid centre; weighted controls span the same climate strata




Supplementary Figure S4 | The national matched design across Chile's aridity gradient.

(a) Matched sample of subcuencas, each filled by long-term aridity (P/PET, the north-south gradient the design conditions on); dammed (treated) basins are outlined in orange and marked with their reservoir (triangles), and each matched-control basin carries a centroid dot sized by its entropy-balancing weight, so the effective control sample (91) is visible against the nominal 244. Dammed basins concentrate in the arid centre and north; weighted controls span the same climate strata. (b) Within-basin placebo geometry for one representative dammed basin: the reservoir (triangle) with its downstream (regulated) gauges below the dam and upstream (unregulated) gauges above it.


Matching removes the siting imbalance on common support

Weighting collapses area/elevation gaps to ~0; pruned units are off-support high-Andes basins

a




b



Supplementary Figure S5 | Covariate balance and common support for the matched design. (a) Love plot of the absolute standardized mean difference between dammed and control basins before (unweighted) and after entropy balancing, for the two balanced covariates (log basin area, mean elevation) and two untargeted diagnostics (latitude, log aridity); the dashed line is the conventional 0.1 balance threshold. Weighting collapses the large siting imbalances to ~0, leaving a small residual aridity difference (0.17) that the doubly-robust outcome model absorbs, and reduces the effective control sample from 244 to 91. (b) Climate-stratum common support: mean elevation of control (grey) and dammed (orange, retained) basins within each Köppen main group; the three

pruned dammed basins (x) fall in the cold (D) and polar (E) strata, whose controls lie far below them, so they are off-support and excluded rather than matched by extrapolation.


Water rights: no induced demand survives once siting is matched out



Supplementary Figure S6 | Water rights show no induced demand once siting is matched out. (a) Event study of cumulative consumptive water-rights density (dammed $\times$ year, ref. 2009; rights per 100 km²): the dammed-versus-control gap appears to widen through the megadrought, the naive induced-demand signal. (b) For both the rights count and the winsorized allocated volume, the expansion ATT (2005-2024; orange point, dark analytic 95% CI) is positive and its analytic interval excludes zero, yet it lies squarely inside the 95% band of the permutation null (grey), what random treatment assignment within climate $\times$ aridity strata already produces given the dammed basins' markedly higher aridity, so it is null under randomization inference ($p_{\text{perm}} = 0.52$ for counts and 0.31 for the winsorized volume, both estimated on the full matched sample with the 99.5th-percentile cap under unit-level stratified permutation; the threshold sensitivity is Supplementary Table S8, and the aridity-overlap subset with its spatial-block inference, a different sample and permutation scheme, is Supplementary Table S9). A conventional matched analysis of this registry would report induced demand; the design-based inference reveals it as siting.


The positive whole-basin ET effect is a barren-land artifact

Masking barren pixels (vegetated cells only) removes the pre-trend failure and the effect



Supplementary Figure S7 | The positive whole-basin ET effect is a barren-land artifact.

Event studies of the dammed-versus-control evapotranspiration response (dammed x year, ref. 2009) for (a) whole-basin ET over all pixels, which shows strongly non-parallel pre-trends (pre-trend slope $p = 3 \times 10^{-5}$) and an apparent post-2010 divergence, and (b) ET restricted to vegetated irrigated (orchard) cells, i.e. with barren and salar pixels masked out. Masking the barren land removes both the pre-trend failure ($p = 0.40$) and the positive effect (the forcing-conditioned slope gap moves from +0.048 to -0.020), confirming that the exclusion of whole-basin ET from the main results is mechanistic (MOD16 ET is contaminated over sparsely vegetated barren surfaces), not a post-hoc drop of an inconvenient result.



Supplementary Figure S8 | Baseline aridity: non-parametric common support. Densities of long-term baseline aridity, $\log(P/PET)$, for control basins (grey area, unweighted; dashed line, entropy-weighted) and dammed basins (orange), with per-basin rug marks (top, dammed; bottom, control). The dammed basins are systematically more arid than the unweighted control mean, the residual imbalance (SMD 0.17) that the doubly-robust outcome model adjusts parametrically; this plot shows the adjustment is interpolation, not extrapolation. All 21 dammed basins lie inside the control range and inside the control 5th–95th percentile band (shaded), and the control density is strictly positive across the entire treated support, so no treated unit relies on out-of-support extrapolation. The hard-balanced sensitivity (Supplementary Table S17), which removes the residual imbalance non-parametrically at the cost of an effective control sample of 19, reaches the same null, so the conclusion does not rest on the parametric adjustment.

Supplementary Tables

Supplementary Table S1 | Equivalence / minimum-detectable-effect bounds on the headline nulls. Observed slope gap, MDE (smallest effect detectable at 80% power from the permutation-null SD), the 90% (TOST) interval, and whether it lies within the negligible region ($\pm 25\%$ of baseline). “baseline” is the untreated (control) SPEI→SSI / SPEI→impact transmission slope from the fixed-effects model. No outcome is formally equivalent to zero; the streamflow design rules out attenuations larger than $\sim 46\text{--}58\%$ of baseline transmission.

Outcome	Observed	MDE	90% CI	Equiv. to 0?	p_perm
streamflow SSI (ITT)	-0.183	0.268	[-0.34, -0.0253]	no	0.39
streamflow SSI (downstream)	-0.165	0.312	[-0.348, 0.0186]	no	0.48
irrigated area (DiD)	0.001	0.010	[-0.00506, 0.00656]	no	0.90
orchard ET (DiD)	-0.020	0.046	[-0.0463, 0.00718]	no	0.25

Supplementary Table S2 | Upstream/downstream placebo comparability. Upstream gauges are higher than downstream, and gauge elevation predicts the transmission slope among undammed controls (-0.21 per $+1000$ m), yet upstream and downstream treated gauges have nearly identical raw transmission slopes, so the equal up/down attenuation is not an elevation artifact.

quantity	value	detail
upstream gauges (elev, mean m)	1003.000	n=28; mean transmission 0.565
downstream gauges (elev, mean m)	459.000	n=40; mean transmission 0.599
control transmission slope vs elevation (per $+1000$ m)	- 0.212	p = 0.00 (n=84 control gauges); elevation predicts transmission, so the up/down gap is bounded with this sensitivity rather than assumed away

Supplementary Table S3 | Positive control on the streamflow estimator. A known buffering slope (`beta_injected`; negative = buffering) is added to treated downstream units; the estimator recovers it (recovered = baseline $-0.17 + \text{beta}$), the no-effect case is not flagged, and randomization inference detects the effect once it exceeds the minimum detectable effect.

beta_injected	recovered	p_perm	detected
0.00	-0.165	0.480	FALSE
-0.15	-0.315	0.135	FALSE
-0.30	-0.465	0.014	TRUE
-0.45	-0.615	0.001	TRUE

Supplementary Table S4 | Siting-confound decomposition ladder (numeric). The dammed-versus-control transmission-slope gap at each rung of Supplementary Figure S3, for streamflow SSI (ITT) and the two difference-in-differences proxy outcomes, plus the within-basin upstream placebo. `perm_p` is the design’s randomization-inference p for the design rung and the placebo.

Outcome	Rung	Slope gap	SE	p_perm
Streamflow SSI (SPEI->SSI slope)	naive (no match, no FE)	-0.197	0.090	-
Streamflow SSI (SPEI->SSI slope)	unit+time FE (unweighted)	-0.165	0.075	-
Streamflow SSI (SPEI->SSI slope)	design (ebal + FE)	-0.183	0.083	0.39
Streamflow SSI (SPEI->SSI slope)	design + within region	-0.050	0.070	-
Irrigated area (DiD)	naive (no match, no FE)	0.001	0.003	-
Irrigated area (DiD)	unit+time FE (unweighted)	0.002	0.002	-
Irrigated area (DiD)	design (ebal + FE)	0.001	0.002	0.91
Irrigated area (DiD)	design + within region	0.001	0.002	-
Orchard ET (DiD)	naive (no match, no FE)	-0.017	0.013	-
Orchard ET (DiD)	unit+time FE (unweighted)	-0.015	0.013	-
Orchard ET (DiD)	design (ebal + FE)	-0.020	0.013	0.24
Orchard ET (DiD)	design + within region	-0.020	0.012	-
Streamflow SSI (SPEI->SSI slope)	upstream placebo (unregulated)	-0.201	0.090	0.48

Supplementary Table S5 | Storage-band trends and megadrought level shift. Top: per-year trend slope of the cross-reservoir peak, trough, and peak-trough amplitude (percent of capacity per year; reservoir fixed effects, cluster-robust). Peak and trough both decline while the amplitude trend is not distinguishable from zero, consistent with a supply-side level decline rather than a loss of refill capacity. Bottom: because the storage dynamics are non-linear over a window that begins just before the megadrought, the linear slope is reported as a summary rate only, and the decline is additionally characterized as a within-reservoir level shift at the 2010 onset (pre versus post means with reservoir fixed effects): peak and trough step down significantly while the amplitude shift is indistinguishable from zero.

Component	Trend (%/yr)	95% CI	p
peak	-1.281	[-2.21, -0.349]	0.013
trough	-1.220	[-1.72, -0.72]	0.000
amplitude	-0.061	[-1.13, 1.01]	0.912

Component	Pre-2010 mean	Post-2010 mean	Shift	95% CI	p
peak	0.867	0.665	-0.210	[-0.364, -0.0555]	0.01400
trough	0.395	0.248	-0.150	[-0.227, -0.0734]	0.00081

Component	Pre-2010 mean	Post-2010 mean	Shift	95% CI	p
amplitude	0.473	0.417	-0.059	[-0.217, 0.0987]	0.47000

Supplementary Table S6 | Aridity-overlap sensitivity. Key results re-estimated on the sub-sample of basins whose baseline log-aridity lies in the treated/control overlap band (all 21 treated basins and 158 of 244 controls fall inside it), to show the parametric aridity adjustment is not masking a true effect where common support is genuine. The streamflow transmission-slope gap is essentially unchanged, and the water-rights expansion ATT shrinks and now straddles zero.

Quantity	Full sample	Aridity overlap	Overlap 95% CI
Streamflow SPEI->SSI slope gap	-0.183	-0.174	-
Water-rights expansion ATT	20.036	11.880	[-5.62, 29.4]

Supplementary Table S7 | Fully non-parametric matching on the aridity-overlap subset.

Every cross-sectional outcome re-estimated after re-fitting entropy balancing *within* the treated/control aridity-overlap band and using the weighting-only estimator (ebal-weighted difference in means, no outcome-model aridity term), so the estimate performs no parametric extrapolation across the aridity regime gap. Cropland-area expansion, orchard-ET buffering and the water-rights count remain null. The winsorized water-rights volume is the only outcome whose stratified-permutation p falls below 0.05 here; Supplementary Table S9 shows this reflects spatial autocorrelation, not a reservoir effect. The **perm p** column is the unit-level stratified randomization-inference p (Köppen $\times$ aridity tercile).

Outcome	ATT	95% CI (HC3)	p_perm	n_t	n_c
Cropland-area expansion (ha/km ²)	-0.5860	[-2.08, 0.905]	0.328	21	158
Orchard ET buffering (dlogET/dSPEI)	-0.0182	[-0.0427, 0.00631]	0.291	18	30
Water-rights count (per 100 km ²)	12.0000	[-4.36, 28.4]	0.420	21	136
Water-rights volume (l/s per km ²)	1.8700	[0.164, 3.58]	0.005	21	136

Supplementary Table S8 | Sensitivity of the water-rights volume null to the winsorization threshold.

The winsorized allocated-volume expansion ATT and its randomization-inference p, with individual consumptive flows capped at the 95th, 99th, 99.5th (the value used in the main text) and 99.9th percentiles of the consumptive-flow distribution. The analytic interval excludes zero at every threshold, but the effect is not significant under randomization inference at any threshold (permutation p from 0.39 to 0.21), so the volume null is not an artefact of the winsorization choice. The cap is also physically validated (second table): the 99.5th-percentile cap corresponds to 0.3 m³ s⁻¹, and in every gauged basin the largest retained right is below the basin's maximum observed monthly mean streamflow (median ratio 0.002), so no physically impossible values survive the cap; the uncapped registry maximum, by contrast, exceeds 2×10^8 l/s.

Winsor pctl	Volume ATT	95% CI	p_perm
95.0	1.837	[0.639, 3.03]	0.389
99.0	2.395	[0.956, 3.83]	0.343
99.5	2.539	[1.07, 4.01]	0.310
99.9	3.211	[1.56, 4.86]	0.209

Quantity	Value	Detail
winsorization cap (99.5th pct of consumptive flows)	287.000 l/s (= 0.3 m ³ /s)	uncapped registry maximum 2.08e+08 l/s
gauged basins where the largest retained right $\leq$ max observed monthly flow	77.000	of 77 gauged basins (100%)
median ratio: largest retained right / max observed flow	0.002	IQR [0.000, 0.004]

Supplementary Table S9 | Spatial autocorrelation and spatially-restricted inference.

Top: Moran's I of each outcome's ATT-model residuals (great-circle distances, five-nearest-neighbour

row-standardized weights, 999 residual permutations). The water-rights outcomes and cropland area show significant positive spatial autocorrelation, which the unit-level stratified permutation does not preserve. Bottom: for the water-rights outcomes, the ATT with its spatially-naive (HC3) interval and stratified-permutation p alongside the cuenca-block-clustered interval and the spatial-block permutation p (treatment swapped at the cuenca level). The apparent water-rights volume effect on the overlap subset (stratified p = 0.005) is not significant once spatial dependence is respected (spatial-block p = 0.077; cuenca-clustered interval spans zero), and the count is null under both.

Outcome	n	Moran's I	Expected	p
Cropland-area expansion	265	0.181	-0.004	0.001
Water-rights count	216	0.599	-0.005	0.001
Water-rights volume	216	0.388	-0.005	0.001
Orchard ET buffering	49	0.069	-0.021	0.198

Outcome	ATT	95% CI (HC3)	p (stratified)	95% CI (cuenca)	p (spatial block)
Water-rights volume (l/s per km2)	1.87	[0.16, 3.58]	0.005	[-0.33, 4.07]	0.077
Water-rights count (per 100 km2)	12.00	[-4.36, 28.36]	0.420	[-2.70, 26.70]	0.581

Supplementary Table S10 | Groundwater-substitution bound from the DGA well hydrograph network. Differential groundwater drawdown between dammed and matched-control basins over the megadrought (2010–2021). Each well’s depletion rate is the OLS slope of Hampel-filtered depth to water on time (m yr^{-1} , positive = falling water table); per-basin rates are aggregated by the outlier-robust median (primary) and, for robustness, the mean. **dammed** / **control** are the weighted group-mean rates; **att_wo** is the entropy-balancing weighting-only ATT (with 95% CI and stratified randomization-inference **perm_p**) and **att_dr** the aridity-adjusted doubly-robust ATT. Coverage is 26 matched basins (13 dammed, 13 control, 213 wells). The differential is not significant in either specification, and the depletion residuals carry no spatial autocorrelation (**moran_I**, **moran_p**), so dammed basins show no faster groundwater drawdown of the kind a substitution confound would require. One scope limitation applies: the DGA monitoring network is not designed around agriculture, so monitored wells need not sit inside the high-value orchard footprints where compensatory pumping concentrates, nor sample the same aquifer strata. The basin-median rates therefore bound basin-scale substitution; field-scale pumping localized outside the monitored network remains a residual caveat.

Spec	Dammed (m/yr)	Control (m/yr)	ATT (WO)	95% CI	p_perm	ATT (DR)	Moran I	Moran p	wells
median (primary)	0.098	0.079	0.089	[-0.17, 0.35]	0.681	0.028	- 0.069	0.752	213
mean (robustness)	0.289	0.076	0.283	[-0.11, 0.67]	0.328	0.486	- 0.014	0.726	213

Supplementary Table S11 | Reservoir-use heterogeneity: irrigation-only robustness.

The treated sample is dominated by irrigation reservoirs (18 of the 21 matched treated basins hold one; three are purely hydropower or potable supply). The headline streamflow SPEI-to-SSI slope gap (**treat:SPEI**, < 0 = apparent buffering) is re-estimated on all treated basins and on the irrigation-only subset (non-irrigation treated basins dropped, controls unchanged), with the design's stratified randomization-inference p. The slope gap and its permutation-null status are unchanged, so the null is not an artefact of pooling hydropower with irrigation reservoirs; **n_treat** columns give the treated basins with streamflow gauges in each case.

Quantity	All treated	p (all)	Irrigation only	p (irrig)	n_treat (all)	n_treat (irrig)
Streamflow slope gap (ITT)	-0.183	0.388	-0.192	0.432	17	16
Streamflow slope gap (upstream placebo)	-0.201	0.482	-0.221	0.456	15	14

Supplementary Table S12 | Snowmelt-confound tests on the upstream/downstream placebo, run on the differential estimand. Upstream reaches are snow-dominated Andean headwaters, so snowmelt storage and release could in principle shape the treated-versus-control transmission differential up- versus downstream independently of regulation. All snow terms use ERA5-Land snow-water-equivalent (SWE) standardized across units, since the product underestimates absolute Andean snowpack. Four tests bound the concern: a forcing-by-snowpack adjustment leaves the up/down equivalence intact while resolving a detectable cross-sectional snow term upstream, so the product is not blind to snow; the equivalence persists among low-snow units (below-median peak SWE); the apparent buffering does not strengthen in high-snow years (snow-year modulation); and the treated-versus-control differential is statistically indistinguishable between the early and late halves of the panel (period interaction $p = 0.46$ upstream, 0.52 downstream), providing no evidence of a snow buffer degrading in parallel with the megadrought. Gauge-level rows cap any snow-induced up-minus-down slope difference at 0.007 against the observed 0.04 ; these rows are restricted to placebo-basin gauges with SWE coverage, so gauge counts differ slightly from Supplementary Table S2. Snowmelt is therefore not the driver of the placebo result.

Quantity	ValueDetail
upstream placebo: treat x SPEI, unadjusted	- SE 0.0904, $p = 0.030$; 15 treated / 44 control units 0.2014
upstream placebo: treat x SPEI, snow-adjusted (spei x peak-SWE z)	- SE 0.0899, $p = 0.070$; snow term 0.0922 (SE 0.0360, $p = 0.165$) 0.013
upstream placebo: treat x SPEI, low-snow units (below-median peak SWE)	- SE 0.0791, $p = 0.304$; 3 treated / 26 control units; peak SWE 0.0829 0.0000-0.0012 m w.e.
upstream placebo: treat x SPEI, early vs late period	- early (≤ 2010) -0.3207 (SE 0.1180, $p = 0.009$); late (> 2010) 0.3207 0.2654 (SE 0.1795, $p = 0.145$); period change (treat x SPEI x late) 0.1034 (SE 0.1394, $p = 0.461$)
upstream placebo: treat x SPEI x annual-SWE anomaly (snow-year modulation)	- SE 0.0472, $p = 0.510$; base spei x snow-year 0.0521 (SE 0.0313) 0.0374, $p = 0.169$
downstream: treat x SPEI, unadjusted	- SE 0.0851, $p = 0.058$; 15 treated / 44 control units 0.1646
downstream: treat x SPEI, snow-adjusted (spei x peak-SWE z)	- SE 0.0921, $p = 0.131$; snow term 0.0296 (SE 0.0262, $p = 0.141$) 0.264
downstream: treat x SPEI, low-snow units (below-median peak SWE)	- SE 0.1207, $p = 0.509$; 6 treated / 23 control units; peak SWE 0.0808 0.0000-0.0005 m w.e.
downstream: treat x SPEI, early vs late period	- early (≤ 2010) -0.2060 (SE 0.1276, $p = 0.112$); late (> 2010) 0.2060 0.3216 (SE 0.1851, $p = 0.088$); period change (treat x SPEI x late) -0.1047 (SE 0.1611, $p = 0.518$)
downstream: treat x SPEI x annual-SWE anomaly (snow-year modulation)	- SE 0.0809, $p = 0.207$; base spei x snow-year 0.0345 (SE 0.1031) 0.0328, $p = 0.297$

Quantity	ValueDetail
gauges: up vs down peak SWE (m w.e., ERA5-Land)	0.0682down 0.0482 (gap 0.0200 m = 0.13 SD); n_up=28, n_down=32
control gauges: transmission slope per SD peak SWE	0.0006SE 0.0282, 95% CI [-0.0553, 0.0553], p = 0.999 (n=82); net of elevation p = 0.808
conservative bound: snow-induced up-down slope difference	0.0075SWE gap 0.13 SD x CI-upper sensitivity 0.0553; observed differential gap -0.04

Supplementary Table S13 | Pre-drought aridity-window sensitivity of the matching.

The 1991–2020 WMO normal used for baseline aridity overlaps eleven megadrought years, so the covariate could in principle absorb post-treatment variation. Recomputing unit aridity on the strictly pre-drought 1991–2009 window leaves the design unchanged: unit aridity is essentially identical across windows (Pearson and Spearman rows), 99% of matched units keep the aridity tercile used as a permutation stratum, and the megadrought change in P/PET (2010–2020 minus 1991–2009) does not differ significantly between treated and weighted control basins and is an order of magnitude smaller than the treated-control aridity gap the matching addresses. The matching therefore does not lean on drought-period or reservoir-influenced variation.

Quantity	Value	Detail
unit aridity, pre-drought (1991-2009) vs full (1991-2020) window: Pearson r	0.9982	Spearman rho 0.9993 (n=265 matched units)
aridity-tercile agreement between windows (permutation strata)	0.9920	share of matched units assigned the same tercile
median P/PET, treated: pre-drought vs full window	0.2730	full 0.227; control pre 0.798, full 0.680
megadrought change in P/PET (2010-2020 minus 1991-2009), treated minus control	- 0.0441	SE 0.0265, p = 0.10 (ebal-weighted); small against the ~0.5 treated-control aridity gap

Supplementary Table S14 | Minimum detectable effect for the cropland pre-trend. The flat pre-2010 differential cropland trend (Fig. 4b) is held to the same equivalence standard as the headline nulls: the minimum detectable pre-window trend slope is built from the permutation-null SD (80% power, two-sided alpha 0.05) and cumulated over the 2000–2009 pre-window as a share of the weighted control cropland level. The design would detect a pre-existing divergence accumulating to about half of the control cropland level by 2009; the observed differential trend is about a seventh of the minimum detectable slope.

Quantity	Value	Detail
pre-2010 differential cropland trend (treat:year, basin fraction per yr)	- 0.000195	perm p = 0.72 (1000 perms); window 2000-2009
pre-trend MDE (80% power, alpha = 0.05, basin fraction per yr)	0.001400	permutation-null SD 0.0005
cumulative detectable pre-2010 divergence (% of control cropland level)	51.300000	MDE x 9 yr against control weighted mean fraction 0.0245

Supplementary Table S15 | Inter-basin spillover and administrative-demand checks.

Top rows: matched treated and effective-control basins occupy entirely disjoint DGA cuencas, so no matched pair shares a river network and the within-cuenca contamination rule already removes the direct transfer pathway. The residual transfer bias is bounded in both directions: a transfer moving water from a dammed basin to a control raises the control's availability and shrinks the estimated gap, biasing toward the null we report, while imports sustaining a treated basin during peak drought would flatten the treated transmission slope where deliveries land, downstream, and so mimic regulation-specific buffering; the absence of any downstream-specific flattening in the placebo bounds this direction too, though a vulnerability signal partially offset by imports cannot be fully excluded without an explicit transfer-network model. Bottom rows: the water-rights registry was not closed during the megadrought, and the closure that would matter, a freeze specific to dammed basins, is absent: treated basins entered the drought with a higher allocated stock per unit area than their weighted controls, accrued new rights about 2.7 times faster through the megadrought, and none went without new grants, while 29% of weighted controls did. The accrual null is therefore a genuine null on new state grants, not the forced silence of saturated registries; what administrative closure does imply is that grant accrual under-represents total demand, because reallocation in closed basins proceeds through private market transfers the grant registry does not record (main text).

Metric	Value	Detail
treated cuencas	14.00	distinct DGA cuencas holding ≥ 1 treated subcuenca
effective-control cuencas	70.00	distinct DGA cuencas holding ≥ 1 effective (w>0) control
cuencas shared by a matched pair	0.00	contamination rule removes same-cuenca controls $\Rightarrow$ no matched pair shares a river network
new consumptive rights per megadrought year	4687.00	range 2539-7352 over 2010-2021; registry not closed
basins granted new rights per year	194.00	of 265 matched basins
pre-drought (2009) rights stock, treated (per 100 km ² , weighted)	15.50	control 6.2; higher stock over lower renewable supply $\Rightarrow$ treated utilization at least as high
pre-drought (2009) winsorized volume stock, treated (l/s per km ² , weighted)	1.84	control 0.65
megadrought accrual rate, treated (rights per 100 km ² per yr, weighted)	1.72	control 0.63; treated accrual not slower $\Rightarrow$ no differential administrative closure
treated basins with zero new rights over the megadrought (share)	0.00	control share 0.29

Supplementary Table S16 | Physical-volume translation of the equivalence margin.

The $\pm 25\%$ equivalence margin corresponds at the worst observed deficit to a perturbation of 0.29 SSI units (Methods). Translated through each treated downstream gauge's within-gauge linear sensitivity of the 12-month rolling mean flow to its SSI, that perturbation corresponds to the volumes below, expressed per 12-month season and as a share of mean annual flow. The linear sensitivity is an average rather than a dry-tail mapping and therefore overstates the worst-drought volume; even so, the median is of order tens of cubic hectometres, a volume that is material for local water management. This is why the margin is anchored to the operational drought-severity classification rather than to volumetric negligibility, why the threshold-free minimum detectable effect is reported alongside, and why the causal weight of the null rests on the within-basin placebo rather than on the equivalence verdict.

Quantity	Value	Detail
equivalence-margin volume, treated downstream gauges (hm ³ per 12-month season)	58.0	median across 40 gauges; IQR [16.3, 137.2]; flow-SSI sensitivity x 0.29 SSI x 31.5; linear (average) sensitivity overstates the dry-tail volume
equivalence-margin volume as % of mean annual flow	10.7	IQR [6.9, 14.0]%

Supplementary Table S17 | Hard-balanced aridity sensitivity. The primary design deliberately leaves a residual aridity imbalance (SMD 0.17) to the doubly-robust outcome model because hard-balancing aridity collapses the effective control sample from 91 to about 19. Here the strict alternative is reported: entropy balancing that targets log baseline aridity alongside log area and elevation within Köppen group, achieving SMD 0.000 with no parametric extrapolation across the hydroclimatic regime gap. Every headline result is unchanged or strengthened: the streamflow slope gap and the upstream placebo collapse essentially to zero, and the cropland and water-rights outcomes remain null. That even the *apparent* buffering vanishes when aridity is fully balanced is the direct signature of a siting confound rather than a reservoir effect. The weight concentration this balance requires, and a leave-one-out sensitivity showing the point-estimate vanishing depends on the single most-weighted control basin, are disclosed in Supplementary Table S45; the hard balance is corroborating, not load-bearing.

Quantity	Value	Detail
effective control sample size (aridity hard-balanced)	18.900000	log-aridity SMD after balancing 0.000 (residual 0.17 in the primary design)
streamflow treat x SPEI, ITT	- 0.023000	SE 0.160, perm p = 0.930 (1000 perms); primary design -0.183
streamflow treat x SPEI, upstream placebo	- 0.047000	SE 0.154, perm p = 0.882; primary design -0.20
cropland-area DiD slope gap	0.000508	SE 0.0023, perm p = 0.855; primary design perm p = 0.91
water-rights expansion ATT (weighting-only, per 100 km ²)	- 4.900000	95% CI [-36.9, 27.1], perm p = 0.855

Supplementary Table S18 | Rank-based tests on raw, uncapped water-rights volumes.

Rights volumes are extremely concentrated and carry known data errors, which is why the level analysis winsorizes; rank-based statistics are immune to caps, outliers, and any monotone error in extreme caudal entries. The per-right volume distribution is nearly identical in dammed and control basins at the median, so the count density does not systematically understate demand in dammed basins. The naive basin-level rank gap in raw added volume reproduces the siting gap, as with the count; the design contrast (ebal-weighted mean-rank difference) appears significant under unit-level permutation, which is anti-conservative given the strong spatial autocorrelation of water-rights outcomes (Supplementary Table S9), and is not significant under the governing cuenca-block permutation, matching the winsorized-volume verdict.

Quantity	ValueDetail
volume per right, raw (median l/s): dammed vs control	2.000control 1.8; Wilcoxon rank-sum $p = 0.00$ ($n = 77,287$ rights)
raw volume added 2005-2024, uncapped (median l/s per km ²): dammed vs control	2.310control 0.17; naive Wilcoxon $p = 0.00$ (265 basins) reproduces the siting gap the design removes
design rank contrast on raw volume (mean-rank difference, ebal-weighted)	84.700unit-level stratified permutation $p = 0.001$, anti-conservative under the spatial autocorrelation of Supp Table S9
design rank contrast under spatially-restricted (cuenca-block) permutation	84.700cuenca-block permutation $p = 0.145$ (1999 perms), the governing inference for water-rights outcomes; rank scale immune to caudal errors and caps
share of total raw volume held by the largest 1% of rights	0.997raw registry (data errors included); why the level analysis winsorizes

Supplementary Table S19 | Volumetric-scale streamflow checks. SSI standardizes each gauge on its own, for regulated gauges post-dam, flow distribution, so the transmission slope tests distribution-relative propagation and would divide out any dam-induced compression of flow variance. Two direct checks show the streamflow null is not a standardization artifact. First, the buffering models and the upstream placebo are re-estimated on log 12-month mean flow, a volumetric semi-elasticity (percent of flow per SPEI unit) that no per-gauge normalization can compress: all slope gaps remain permutation-null, with point estimates slightly positive (toward steeper, not buffered, transmission) and the upstream placebo at least as large as downstream, reproducing the siting pattern. Second, the compression premise itself: deseasonalized log-flow variability at regulated downstream gauges is statistically indistinguishable from unregulated upstream gauges (Wilcoxon $p = 0.43$), so there is no differential variance compression for the standardization to divide out at the monthly-to-annual scale our design resolves.

Quantity	Value Detail
log-flow semi-elasticity gap, ITT (treat x SPEI)	0.0743SE 0.0720, perm $p = 0.562$; baseline $d \log Q / d \text{ SPEI}$ 0.139 (% flow per SPEI unit)
log-flow semi-elasticity gap, downstream (treat x SPEI)	0.0401SE 0.0855, perm $p = 0.782$; baseline $d \log Q / d \text{ SPEI}$ 0.099 (% flow per SPEI unit)
log-flow semi-elasticity gap, upstream placebo (treat x SPEI)	0.0952SE 0.0682, perm $p = 0.422$; baseline $d \log Q / d \text{ SPEI}$ 0.168 (% flow per SPEI unit)
deseasonalized log-flow SD by gauge class (relative variability)	0.8270downstream 0.827 (n=61) vs upstream 0.902 (n=35) vs control 0.579 (n=121); up-down Wilcoxon $p = 0.43$

Supplementary Table S20 | Metric and functional-form checks. Top: the baseline SPEI-to-SSI transmission slope among undammed controls varies smoothly and monotonically across aridity terciles, with no regime break for the linear log-aridity adjustment to misfit. Middle: the control elevation sensitivity of transmission shows no detectable nonlinearity across the ~2000 m snowline (piecewise interaction $p = 0.25$); the snow regime itself is tested directly on the placebo estimand in Supplementary Table S12. Bottom: the cropland null is metric-invariant: re-fit on log cropland share (relative rather than absolute expansion), the forcing-DiD slope gap remains permutation-null.

Quantity	Value	Detail
control baseline SPEI->SSI slope by aridity tercile (arid / mid / humid)	0.6470 mid 0.699, humid 0.743;	smooth monotone variation, no regime break
control elevation sensitivity of transmission, below vs above the ~2000 m snowline	- 0.0780 (n=84 control gauges)	above 0.212 per km; interaction $p = 0.25$
relative cropland expansion: forcing-DiD slope gap on log cropland share	0.0473 SE 0.046, perm $p = 0.369$	(140 units with nonzero cropland); metric-invariant null

Supplementary Table S21 | Evapotranspiration outcomes (relegated from the main text). The ET pathway is reported here with its caveats and carries no evidential weight in the paper’s conclusions. Whole-basin ET fails its event-study pre-trend test (differential pre-2010 trend p approximately $3e-5$) and is confounded by residual aridity over barren land; masking barren pixels removes both the pre-trend failure and the positive effect (Supplementary Fig. S7), but the resulting orchard stratum was defined post hoc and relies on MOD16, whose documented underestimation of ET over irrigated land biases its slope gap toward zero, disqualifying a null as evidence in a null-finding design. A thermal ET product (Landsat SSEBop/METRIC or ECOSTRESS class) is the way to test the consumptive-use pathway properly and is flagged as future work.

Group	Outcome	Estimand	Estimate	95% CI	p
Cross-sectional matched ATT	Reservoir ET buffering	d log ET / d SPEI	-0.014	[-0.0646, 0.0369]	-
Forcing-interacted DiD (slope-gap)	Whole-basin ET	differential deficit->impact slope	0.048	[0.0255, 0.0703]	0.11
Forcing-interacted DiD (slope-gap)	Orchard ET	differential deficit->impact slope	-0.020	[-0.0448, 0.00565]	0.24

Supplementary Table S22 | Geomorphological (alluvial) confound check on the placebo.

The up/down contrast coincides with the bedrock-to-alluvial transition, and elevation is only a proxy for it, so a direct terrain metric is extracted: local relief, the maximum minus minimum SRTM elevation within 2.5 km of each gauge, a standard discriminator of bedrock canyons (high relief) from alluvial valleys (low relief). The up/down relief gap is real (about 0.7 SD), but relief does not predict the transmission slope among undammed controls ($p = 0.41$; marginal at $p = 0.08$ net of elevation), a relief-by-forcing adjustment leaves both placebo coefficients essentially unchanged, and among low-relief (alluvial) units neither the upstream nor the downstream contrast shows any buffering (both positive and null, on few treated units). Alluvial storage therefore neither manufactures the upstream attenuation nor masks a downstream reservoir effect; relief at 2.5 km cannot resolve aquifer storage volumes themselves, which remain a residual caveat.

Quantity	Value	Detail
gauges: up vs down local relief (m, 2.5 km radius)	796.0000	down 502 (gap 294 m = 0.72 SD); n_up=28, n_down=32
control gauges: transmission slope per SD relief	0.0263	SE 0.0320, $p = 0.413$ (n=84); net of elevation $p = 0.084$
upstream placebo: treat x SPEI, unadjusted	- 0.2014	SE 0.0904, $p = 0.030$
upstream placebo: treat x SPEI, relief-adjusted (spei x relief z)	- 0.1968	SE 0.0912, $p = 0.035$; relief term -0.0132 (SE 0.0403, $p = 0.745$)
upstream placebo: treat x SPEI, low-relief (alluvial) units	0.0602	SE 0.1664, $p = 0.721$; 3 treated / 26 control units below median relief
downstream: treat x SPEI, unadjusted	- 0.1646	SE 0.0851, $p = 0.058$
downstream: treat x SPEI, relief-adjusted (spei x relief z)	- 0.2035	SE 0.1052, $p = 0.058$; relief term -0.0483 (SE 0.0439, $p = 0.276$)
downstream: treat x SPEI, low-relief (alluvial) units	0.1250	SE 0.1268, $p = 0.333$; 7 treated / 22 control units below median relief

Supplementary Table S23 | Carryover capacity of the treated reservoirs. Whether the 12-month null tests a function the dams were designed to perform is answerable from the project’s own data. The median treated reservoir stores about 0.3 years of mean annual downstream flow, seasonal-scale by design, but seven of the 17 reservoirs with a gauged downstream unit hold at least half a year and four hold more than a full year of flow; and seasonal operation does not empty the fleet, whose pre-2010 annual trough averaged 39% of capacity, so storage persisted across years. Interannual carryover was therefore physically available to a substantial subset, and the multi-year null is not a deterministic consequence of purely seasonal design; the irrigation-only restriction (Supplementary Table S11) additionally shows the result is not diluted by reservoir purpose. Downstream flow is regulated, so the ratios are approximate.

Quantity	Value	Detail
reservoir capacity / mean annual downstream flow (median ratio)	0.29	IQR [0.15, 0.73]; n = 17 reservoirs with a gauged downstream unit; downstream flow is regulated, so ratios are approximate
reservoirs with carryover ratio ≥ 0.5	7.00	of 17; 4 have ratio ≥ 1 (storage exceeding a full year of flow)
pre-2010 mean annual storage trough (fraction of capacity)	0.39	seasonal drawdown does not empty the reservoirs; storage persists across years, so carryover buffering was physically available

Supplementary Table S24 | Early/late megadrought phase split of the streamflow buffering and placebo tests. A single 2005–2024 transmission slope could in principle average away an early-phase buffering signal (storage intact at drought onset) that fails later (storage depleted). Splitting the megadrought at 2015 shows the opposite of a depletion signature. In the early phase (2010–2014), when reservoirs entered near-full (peak 87% of capacity, Supplementary Table S5) and depletion cannot yet explain anything, the downstream and upstream placebo gaps are identical (treat x SPEI = -0.43 in both), so the strong early attenuation is siting, exactly as in the pooled analysis. The attenuation then weakens from early to late phase, but this weakening is present, and is in fact sharpest, in the unregulated upstream placebo itself (shift $+0.36$, $p = 0.022$), which has no storage to deplete; it therefore reflects the changing forcing regime of the prolonged drought, not the exhaustion of reservoir buffering. At no phase does any subset show a downstream-specific effect beyond its upstream placebo. Per-phase p-values are cluster-robust and shown for completeness; the design’s primary inference remains the randomization test on the pooled estimand (main text).

Quantity	Value	Detail
ITT: treat x SPEI, early megadrought (2010-2014)	- 0.4442370	SE 0.1235, $p = 0.001$; 17 treated / 44 control units
ITT: treat x SPEI, late megadrought (2015-2024)	- 0.3408339	SE 0.2611, $p = 0.197$; 17 treated / 43 control units
ITT: early-to-late phase shift (treat x SPEI x late)	0.2879824	SE 0.1544, $p = 0.067$; pooled 2010-2024 panel, year FE absorb the phase level
downstream: treat x SPEI, early megadrought (2010-2014)	- 0.4299787	SE 0.1028, $p = 0.000$; 15 treated / 44 control units
downstream: treat x SPEI, late megadrought (2015-2024)	- 0.4441582	SE 0.3616, $p = 0.224$; 14 treated / 43 control units
downstream: early-to-late phase shift (treat x SPEI x late)	0.2927161	SE 0.1567, $p = 0.067$; pooled 2010-2024 panel, year FE absorb the phase level
upstream placebo: treat x SPEI, early megadrought (2010-2014)	- 0.4301870	SE 0.1589, $p = 0.009$; 15 treated / 44 control units
upstream placebo: treat x SPEI, late megadrought (2015-2024)	- 0.2148096	SE 0.2247, $p = 0.343$; 15 treated / 43 control units
upstream placebo: early-to-late phase shift (treat x SPEI x late)	0.3643440	SE 0.1547, $p = 0.022$; pooled 2010-2024 panel, year FE absorb the phase level

Supplementary Table S25 | Direct test of the upstream/downstream placebo contrast.

The paper’s central falsifier, tested directly rather than asserted from the separate point estimates: $D = \beta_{\text{down}} - \beta_{\text{up}}$, the downstream-minus-upstream difference in the differential (treat $\times$ SPEI) transmission slope, against its permutation null (treatment reassigned at the unit level within Köppen $\times$ aridity-tercile strata, the same draw applied to both gauge-class panels; Methods). A genuine regulation effect predicts $D < 0$; the observed contrast is +0.037 ($p = 0.72$), so the regulated reaches show no attenuation beyond the unregulated placebo of the same basins.

Quantity	Estimate	Detail
within-basin placebo contrast (downstream - upstream differential slope)	0.037	downstream -0.165, upstream placebo -0.201, $D = +0.037$ (SE_perm 0.098); permutation $p = 0.718$ (2000 perms); a genuine regulation effect would give $D < 0$

Supplementary Table S26 | Within-region residual (siting-ladder rung 4), permutation inference. The design estimator with the baseline SPEI-to-outcome slope allowed to vary by Köppen region, the “regulation estimate” left after the between-region aridity gradient is removed (rung 4 of Supplementary Fig. S3 and Table S4), reported with its own permutation null: the residual is -0.05 ($p = 0.65$).

Quantity	Estimate	Detail
siting-ladder rung 4 (design + within region), permutation p	-0.05	est -0.050, permutation $p = 0.652$ (2000 perms), null SD 0.110

Supplementary Table S27 | Permutation inference for the megadrought phase shift.

The early-to-late change of the buffering slope (treat $\times$ SPEI $\times$ late-phase interaction on the pooled megadrought panel; Methods) under the design’s permutation inference, for the downstream panel, the upstream placebo, and their contrast. Neither phase shift is significant and the down-minus-up contrast is far from it; the largest shift is the upstream placebo’s, which has no storage to deplete, so the phase evolution reflects the changing forcing regime of the prolonged drought rather than buffer exhaustion (companion to the per-phase subsample estimates in Supplementary Table S24).

Quantity	Value	Detail
downstream: early-to-late phase shift (treat $\times$ SPEI $\times$ late), permutation p	0.293	est +0.293, permutation $p = 0.355$ (2000 perms)
upstream placebo: early-to-late phase shift (treat $\times$ SPEI $\times$ late), permutation p	0.364	est +0.364, permutation $p = 0.088$ (2000 perms)
phase-shift contrast (downstream shift - upstream placebo shift)	- 0.072	est -0.072, permutation $p = 0.798$ (2000 perms); a regulation effect would give contrast < 0

Supplementary Table S28 | Non-linear transmission test. A quadratic SPEI term interacted with treatment on the full panel: a buffering effect confined to extreme deficits, as threshold-based release rules would produce, would appear as a nonzero $\text{treat} \times \text{SPEI}^2$ coefficient. It is null.

Quantity	Es- ti- mate	Detail
non-linear transmission test ($\text{treat} \times \text{SPEI}^2$), permutation p	- 0.025	est -0.025, permutation p = 0.654 (2000 perms), null SD 0.042; a dry-tail-only buffering effect would appear as a nonzero interaction

Supplementary Table S29 | Dry-tail re-estimates of the buffering and placebo tests. The transmission models and the up/down contrast re-fit on meteorological drought months only ($\text{SPEI} < -1.0$, moderate drought and drier; and $\text{SPEI} < -1.5$ as a lower-power sensitivity), each under permutation inference. No subset shows a downstream-specific effect beyond its upstream placebo; the < -1.5 rows retain few months per gauge and are reported for completeness.

Quantity	Value	Detail
dry-tail buffering, ITT ($\text{SPEI} < -1.0$)	0.001	est +0.001, permutation p = 1.000 (2000 perms)
dry-tail buffering, downstream ($\text{SPEI} < -1.0$)	- 0.199	est -0.199, permutation p = 0.711 (2000 perms)
dry-tail buffering, upstream placebo ($\text{SPEI} < -1.0$)	0.052	est +0.052, permutation p = 0.939 (2000 perms)
dry-tail placebo contrast down-minus-up ($\text{SPEI} < -1.0$)	- 0.250	D = -0.250, permutation p = 0.487 (2000 perms); regulation would give D < 0
dry-tail buffering, ITT ($\text{SPEI} < -1.5$)	0.584	est +0.584, permutation p = 0.512 (2000 perms)
dry-tail buffering, downstream ($\text{SPEI} < -1.5$)	2.518	est +2.518, permutation p = 0.425 (2000 perms)
dry-tail buffering, upstream placebo ($\text{SPEI} < -1.5$)	- 0.709	est -0.709, permutation p = 0.518 (2000 perms)
dry-tail placebo contrast down-minus-up ($\text{SPEI} < -1.5$)	3.227	D = +3.227, permutation p = 0.453 (2000 perms); regulation would give D < 0

Supplementary Table S30 | Catchment-area comparability of the placebo. The unit-area gap between upstream and downstream gauges is negligible, and log unit area does not predict the transmission slope among undammed control gauges (with a tercile check for non-monotonicity), so catchment size cannot manufacture the upstream attenuation. The unit area is the sub-watershed polygon each gauge sits in; the cumulative-scale confound, which the polygon area does not capture, is tested with gauge discharge in Supplementary Table S33.

Quantity	Value	Detail
upstream gauges (unit area, km ²)	3160.00	n=28; mean transmission 0.565
downstream gauges (unit area, km ²)	3070.00	n=40; mean transmission 0.599
control transmission slope vs log(area), per e-fold	- 0.07	p = 0.06 (n=84 control gauges); area does not predict transmission, so the up/down catchment-size gap cannot manufacture the placebo attenuation
control transmission by area tercile (mean slope)	NA	small 0.865 (n=28) mid 0.891 (n=28) large 0.725 (n=28)

Supplementary Table S31 | Elevation-adjusted placebo. The buffering model re-fit with a SPEI-by-elevation interaction, so the differential slope is net of the elevation-transmission gradient (which alone would make upstream reaches look *more* buffered; main text). The elevation term itself is strongly detectable (`elev_term_p`), yet every treatment effect and the down-minus-up contrast remain null, so the placebo result is not an elevation artifact.

Subset	Raw est.	Adjusted est.	p_perm (adj)	Elev. term p	n perms
ITT	-0.1828730	-0.1952412	0.2998501	0.00043	2000
downstream	-0.1645997	-0.1960080	0.3308346	0.00074	2000
upstream placebo	-0.2014380	-0.2117479	0.3903048	0.00088	2000
contrast down-up	NA	0.0157399	0.8905547	-	2000

Supplementary Table S32 | Gauge-classification robustness of the placebo contrast.

The contrast re-estimated (i) excluding downstream gauges whose DGA hydrographic sub-cuenca is named for a different river or an unregulated tributary than the dam and (ii) excluding downstream gauges within 100 m of the dam elevation, the ones most plausibly on a parallel tributary. The contrast stays null in every scenario, and the downstream estimate moves toward zero rather than toward buffering.

Scenario	n	Down	Up	D	p_D	p
	down					(down)
baseline (all downstream gauges)	61	- 0.1645997	-0.201438	0.0368383	0.723	0.481
excluding cross-river gauge(s)	59	- 0.1606022	-0.201438	0.0408358	0.688	0.494
excluding gauges within 100m of dam	45	- 0.0823384	-0.201438	0.1190996	0.470	0.504

Supplementary Table S33 | Cumulative-scale (discharge-adjusted) placebo. Because larger cumulative catchment scale genuinely flattens transmission among controls (+0.08 per log discharge, $p < 0.001$), the buffering model adds a SPEI-by-log-discharge interaction so the differential slope is net of cumulative scale. Every effect and the contrast remain null, and the scale gap is small in most paired basins (median downstream/upstream discharge ratio 1.1), because downstream gauges sit just below the dams.

Quantity	ValueDetail
scale-adjusted buffering, ITT (SPEI x log discharge)	- raw -0.183; permutation p = 0.502 (1000 perms) 0.093
scale-adjusted buffering, downstream	- raw -0.165; permutation p = 0.504 (1000 perms) 0.102
scale-adjusted buffering, upstream placebo	- raw -0.201; permutation p = 0.678 (1000 perms) 0.072
scale-adjusted placebo contrast (down - up)	- permutation p = 0.737 (1000 perms); regulation would give < 0 0.030
control transmission slope vs log discharge (cumulative scale)	0.078 p = 0.000 (n=84 control gauges); larger cumulative scale flattens transmission, so the contrast must be scale-adjusted
downstream/upstream mean-discharge ratio (paired basins)	1.110 median 1.11, IQR [0.94, 2.26], n=15; 10 of 15 basins have ratio < 2, so downstream cumulative area is NOT vastly larger in most paired basins

Supplementary Table S34 | Paired-placebo contrast power. The minimum detectable downstream-minus-upstream contrast at 80% power (two-sided $\alpha = 0.05$), built from the permutation-null SD: 0.28 SSI units, about 52% of the downstream baseline transmission. The paired test therefore discriminates a siting signature from *large* regulation effects only; moderate downstream-specific buffering is not excluded, and the causal weight rests on the convergence of the elevation-, scale-, misclassification-, and extraction-adjusted contrasts rather than on any single test.

Quantity	ValueDetail
paired-placebo contrast MDE (downstream minus upstream buffering gap)	0.275observed D +0.037 (perm p 0.718), null SD 0.098; MDE 0.275 SSI units = 51.8% of the 0.53 downstream baseline at 80% power, alpha 0.05; a true downstream-specific buffering gap above this is detected, smaller gaps are not excluded (n = 15 paired basins)

Supplementary Table S35 | Storage-capacity (carryover) heterogeneity of the buffering test. The buffering sample (17 treated basins with a usable downstream gauge) split by the reservoir carryover ratio (capacity / mean annual downstream flow, in years): ITT, downstream, and upstream slopes and the down-minus-up contrast D per subgroup, each under permutation inference. The seven carryover reservoirs (ratio ≥ 0.5) show a downstream-specific contrast ($D = -0.23$, $p = 0.014$) that the ten seasonal reservoirs do not ($D = +0.18$, $p = 0.26$). This split is an exploratory, reviewer-prompted subgroup analysis and carries no multiplicity adjustment; its hardening battery is in Supplementary Tables S39 (strict pairing), S40 (denominator sensitivity), and S42 (hydropower exclusion and leave-one-out inference).

Subset	n treat	ITT					Perm. p
		ITT	p	Down	Up	D	
carryover reservoirs (capacity ≥ 0.5 yr)	7	-0.323	0.247	-0.459	-0.227	-0.232	0.014
seasonal reservoirs (capacity < 0.5 yr)	10	-0.078	0.601	0.001	-0.178	0.178	0.258

Supplementary Table S36 | Downstream extraction masking. The up/down placebo compares regulated downstream reaches with unregulated upstream reaches of the same basins, and downstream valleys carry far more irrigation than upstream headwaters; drought-intensified downstream extraction could in principle steepen the downstream transmission slope and mask a real buffer. Rows: the downstream-versus-upstream local cropland gradient within the paired basins (MapBiomass class 18 in 5 km gauge buffers, aggregated per basin-class); the per-gauge control SPEI→SSI transmission slope regressed on local cropland fraction, with and without elevation and cumulative scale; and the masking magnitude predicted by the gradient times each sensitivity. The gradient is real (downstream gauges sit in roughly two and a half times the cropland of upstream gauges; 12 of 15 paired basins downstream-richer), but local extraction does not predict the transmission slope among control gauges, the point estimate carries the sign opposite to the masking mechanism, and even at face value the gradient predicts a slope shift near 0.01 SSI units, two orders of magnitude below the 0.28 contrast MDE.

Quantity	Value	Detail
downstream-gauge local cropland (2007, 5 km)	0.197457	per-basin mean over 15 paired basins; downstream valleys are ~2.5x more cropped than upstream headwaters
upstream-gauge local cropland (2007, 5 km)	0.077435	per-basin mean over 15 paired basins
down-up cropland gap (paired treated gauges)	0.120021	positive = downstream more extracted; 12 of 15 paired basins have down > up
control transmission slope vs local cropland (raw)	- 0.046579	p = 0.749 (n=84 control gauges); local extraction does NOT predict transmission
control transmission slope vs local cropland (+ elevation)	- 0.100265	p = 0.432; sign opposite to the masking mechanism
control transmission slope vs local cropland (+ elevation + discharge)	0.100265	high-extraction reaches are if anything flatter)
predicted masking magnitude (gap x raw sensitivity)	- 0.094822	p = 0.407 (n=84); robust to scale/elevation
predicted masking magnitude (gap x full sensitivity)	0.005590	= -0.0056 SSI/SPEI, far below the 0.28 MDE and the -0.23 carryover contrast
	0.0113807	= -0.0114 SSI/SPEI, ~3% of the 0.28 MDE

Supplementary Table S37 | Placebo net of a SPEI-by-local-cropland interaction. The direct control for downstream extraction: the buffering model adds SPEI x local-cropland, so the differential slope is net of the extraction-driven transmission gradient, for the ITT, downstream, and upstream panels and the down-minus-up contrast, each under the design’s unit-level permutation inference. Adjusting for local extraction leaves the contrast null ($D = -0.03$, $p = 0.83$); the aggregate null is therefore not a masked buffer.

Subset	Estimate	Perm. p	n perms
ITT	-0.2416300	0.1998002	1000
downstream	-0.2441015	0.2437562	1000
upstream placebo	-0.2123475	0.5624376	1000
contrast down-up	-0.0317540	0.8281718	1000

Supplementary Table S38 | Sub-watershed SPEI exposure misclassification. Upstream and downstream gauges within a paired basin sit in the same sub-watershed and receive the same subcuena-average SPEI-12, so each gauge’s true upslope-catchment exposure is misclassified to the extent that catchment differs from the subcuena mean. Rows: per gauge class, the correlation between gauge-point SPEI-12 (a conservative proxy for catchment exposure, since a catchment mean has lower error variance than any single cell) and the subcuena-mean SPEI the model uses, the error SD and variance, the SPEI variance, and the implied transmission-slope attenuation factor $1/(1 + \text{var}(\text{err})/\text{var}(\text{SPEI}))$. Exposure error is small and nearly symmetric (attenuation 0.95 downstream versus 0.93 upstream); the asymmetry biases the placebo contrast by at most about 0.01 units, and in the direction (upstream slightly more attenuated) that works against downstream- specific buffering.

Class	N	r	Error SD	Error var	SPEI var	Attenuation
control	36300	0.9783510	0.1960651	0.0384415	0.8550878	0.9569779
treated_down	18300	0.9751562	0.2116526	0.0447968	0.8683456	0.9509421
treated_up	10500	0.9671283	0.2418160	0.0584750	0.8298157	0.9341714

Supplementary Table S39 | Strictly-paired carryover contrast. The capacity split (7 carryover + 10 seasonal = 17) is on the buffering sample, the treated basins with a usable downstream gauge; the down/up contrast within a subgroup uses each subgroup's available panels, which are only partly paired. Restricting both sides of the carryover contrast to the six basins with usable gauges in both classes leaves it essentially unchanged ($D = -0.234$, permutation $p = 0.004$, versus $D = -0.232$, $p = 0.014$ in the full split), so the carryover buffering does not rest on the partially-paired composition.

Subset	n paired	D	Perm. p	Down	Up	n perms
strictly-paired carryover (n=6)	6	-0.234	0.004	-0.459	-0.225	1000

Supplementary Table S40 | Carryover-ratio denominator sensitivity. The carryover ratio (capacity / mean annual flow) uses regulated downstream flow, which consumption can depress; recomputed with the unregulated upstream-gauge flow, the available inflow proxy, the classification is stable for all seven carryover basins and moves one seasonal basin (Corrales, whose upstream gauge drains only a small headwater) into the carryover group. The heterogeneity result keeps its direction under the upflow-based split (carryover $D < 0$, seasonal $D = 0$) but the carryover contrast weakens to $D = -0.165$ ($p = 0.146$) because the added basin does not buffer; the direction is robust to the ratio’s denominator, the significance is not.

Unit	Reservoir	Capacity (hm3)	Down-flow ratio	Up-flow ratio	Down-flow class	Up-flow class	Flip
210	Conchi	22	0.732	0.899	carryover	carryover	
343	Lautaro	16	0.314	0.311	seasonal	seasonal	
382	Santa Juana	176	1.160	1.403	carryover	carryover	
432	Puclaro	209	0.718	0.617	carryover	carryover	
450	Recoleta	100	2.218	1.075	carryover	carryover	
453	Cogotí	156	2.546	3.176	carryover	carryover	
471	Corrales	50	0.148	2.002	seasonal	carryover	yes
472	El Bato	26	0.294	0.356	seasonal	seasonal	
603	Convento Viejo	237	0.193	0.484	seasonal	seasonal	
732	Colbún	1544	0.658	0.511	carryover	carryover	
733	Digua	225	0.244	0.300	seasonal	seasonal	
735	Bullileo	60	0.033	NA	seasonal	no-up-gauge	yes
811	Coihueco	30	0.039	0.156	seasonal	seasonal	
831	Ralco	1174	0.180	0.380	seasonal	seasonal	
837	Lago Laja	5582	7.522	28.619	carryover	carryover	

Subset	n treat	ITT	ITT p	Down	Up	D	Perm. p
carryover reservoirs (capacity ≥ 0.5 yr)	8	-0.293	0.264	-0.383	-0.218	-0.165	0.146
seasonal reservoirs (capacity < 0.5 yr)	9	-0.077	0.584	0.004	-0.184	0.188	0.272

Supplementary Table S41 | Storage-band step-versus-trend decomposition. The storage decline nested at the 2010 megadrought onset: a level step at 2010 plus a post-2010 drift, for the cross-reservoir peak and trough (reservoir fixed effects, cluster-robust). The peak steps down at the onset (-0.19 of capacity, $p = 0.04$) and shows no further drift ($p = 0.49$); the trough step is not significant ($p = 0.12$) but the trough keeps falling through the megadrought (-0.010 yr^{-1} , $p = 0.007$). The full-window linear trend of Supplementary Table S5 therefore mixes a level shift into an apparent drift: the peak rests at a lower equilibrium while the carryover floor continues to erode.

Component	Quantity	Estimate	SE	p
peak	within-megadrought trend (per yr)	-0.002	0.003	0.4900
trough	within-megadrought trend (per yr)	-0.010	0.003	0.0065
peak	step change at 2010	-0.192	0.088	0.0400
peak	post-2010 drift (per yr)	-0.002	0.003	0.4900
trough	step change at 2010	-0.074	0.046	0.1200
trough	post-2010 drift (per yr)	-0.010	0.003	0.0065

Supplementary Table S42 | Carryover-claim hardening: hydropower exclusion and leave-one-out inference. The carryover class is not a random draw of the fleet: it mixes five irrigation reservoirs (Conchi, Santa Juana, Puclaro, Recoleta, Cogotí) with two large dams whose operation involves hydropower generation (Colbún, registry use irrigation/generation; Lago Laja, operated jointly for irrigation and generation), whose scheduling could smooth flows for reasons unrelated to drought carryover. The first row re-estimates the carryover contrast on the irrigation-only carryover basins (controls unchanged). The remaining rows re-run the contrast dropping one carryover basin at a time, each with its own permutation null, so the stability of the p-value, not only of the point estimate, is visible; with seven treated units a single influential basin could otherwise move the permutation p across conventional thresholds.

Subset	n treat	ITT					Perm. p
		ITT	p	Down	Up	D	
irrigation-only carryover (hydropower excluded, n=5)	5	-	0.246	-	-	-	0.057
		0.402		0.628	0.321	0.306	
leave-one-out: drop Conchi	6	-	0.254	-	-	-	0.032
		0.355		0.478	0.249	0.229	
leave-one-out: drop Santa Juana	6	-	0.317	-	-	-	0.066
		0.258		0.401	0.175	0.226	
leave-one-out: drop Puclaro	6	-	0.336	-	-	-	0.020
		0.281		0.397	0.165	0.232	
leave-one-out: drop Recoleta	6	-	0.270	-	-	-	0.004
		0.344		0.459	0.225	0.234	
leave-one-out: drop Cogotí	6	-	0.278	-	-	-	0.050
		0.340		0.470	0.237	0.233	
leave-one-out: drop Colbún	6	-	0.283	-	-	-	0.022
		0.367		0.549	0.266	0.283	

Subset	n treat	ITT ITT	p	Down	Up	D	Perm. p
leave-one-out: drop Lago Laja	6	-	0.264	-	-	-	0.120
		0.342		0.488	0.269	0.219	

Supplementary Table S43 | Cascade-dam contamination screen of the upstream placebo.

The placebo assumes upstream reaches are physically unregulated, which cascade systems can violate (Ralco above Pangué on the Biobío; the Laguna del Maule works above Colbún; La Laguna above Puclaro on the Elqui): an upstream gauge below another dam would carry real regulation misread as siting. Every upstream gauge is screened against the 1,370-dam national inventory. A *material* dam is DGA-monitored, registry-classified GRANDE, or of wall height at least 15 m (the ICOLD large-dam threshold); dam elevations are SRTM, the paper's classification convention. The within-sub-watershed tier flags gauges below a material dam inside their own treated sub-watershed; the named-cascade tier adds the gauges below the historically documented upstream works on the same river systems (La Laguna 1937 above Puclaro; the Laguna del Maule works 1951 above Colbún; Lagunas del Huasco 1911 above Santa Juana), the true cascades that sit outside the treated sub-watershed; and the same-cuenca tier flags gauges below any material dam anywhere in their DGA cuenca, a conservative bound that necessarily over-flags parallel tributaries. The placebo survives all three exclusions, and the direction of the change is diagnostic: hidden upstream regulation predicts the cleaned upstream slope should steepen toward zero attenuation, and instead it deepens slightly. Under the strict tier-2 exclusion the carryover contrast keeps its direction and grows while its permutation p rises on the thinned upstream panel, consistent with its exploratory label.

Quantity	Value	Detail
upstream gauges screened	35.000	cross 16 treated basins; screened against 445 inventory dams inside treated basins and 54 material dams (monitored / GRANDE / wall ≥ 15 m) in the same cuencas, from the 1,370-dam national inventory
tier 1: below a material dam in the same sub-watershed	3.000	units: 0733, 0831
named-cascade tier: tier 1 plus gauges below documented upstream works	10.000	adds gauges below La Laguna (1937, above Puclaro), the Laguna del Maule works (1951, above Colbun), and Lagunas del Huasco (1911, above Santa Juana); units: 0382, 0432, 0732, 0733, 0831
tier 2: below any material dam in the same cuenca (conservative bound)	20.000	units: 0382, 0432, 0450, 0471, 0542, 0573, 0603, 0732, 0733, 0831; units losing all upstream gauges: 0382, 0432, 0471, 0542, 0573, 0603, 0732, 0733, 0831; includes true cascades (Laguna del Maule 1951 above Colbún; La Laguna 1937 above Puclaro; Lagunas del Huasco 1911 above Santa Juana) and parallel-tributary over-flags
upstream placebo slope, tier-1-cleaned	0.245	- perm p = 0.280; contrast D = +0.081, p = 0.448 (1000 perms); raw upstream -0.20, raw D = +0.04
upstream placebo slope, named-cascade-cleaned	0.204	- perm p = 0.357; contrast D = +0.039, p = 0.792 (1000 perms); n upstream gauges retained 25 of 35
upstream placebo slope, tier-2-cleaned	0.212	- perm p = 0.471; contrast D = +0.048, p = 0.817 (1000 perms); n upstream gauges retained 15 of 35

Quantity	Value	Detail
carryover contrast D, tier-2-cleaned upstream	- 0.359	down -0.459, up -0.100, perm p = 0.127 (n_treat = 7); baseline D = -0.232, $p = 0.014$

Supplementary Table S44 | Capacity-weighted and absolute-volume storage-band trends. The main-text band trends pool reservoirs with equal weight although capacities span three orders of magnitude, so a steep decline in small reservoirs could in principle dominate the pooled percent-of-capacity trend while the large reservoirs holding most of the national buffer stayed stable. Re-fitting with capacity weights shows the opposite: the weighted decline is steeper than the unweighted one, the peak stabilizes after 2010 while the trough keeps falling, and the weighted seasonal amplitude widens ($+0.8 \text{ pp yr}^{-1}$); the flat unweighted amplitude and the widening weighted amplitude both run opposite to the narrowing that sedimentation or a failing refill cycle would produce. The p-values here are analytic and cluster-robust; because capacity weighting concentrates effective weight on a few large reservoirs and so sharpens the small-cluster problem, the governing inference for these coefficients is the wild cluster bootstrap of Supplementary Table S54, under which the weighted trough trend and drift and the amplitude widening survive while the weighted peak trend and the weighted 2010 steps do not. In absolute volume, the coverage-stable subset (22 reservoirs, 98% of fleet capacity) loses $88 \text{ hm}^3 \text{ yr}^{-1}$ of peak and $209 \text{ hm}^3 \text{ yr}^{-1}$ of trough storage over 2005–2024. The registry capacities are static nominal values; the sedimentation reading of the percent-of-capacity decline is bounded in the main text (step timing, unchanged amplitude, co-movement with unregulated control streamflow).

Quantity	Value	Detail
capacity-weighted trend (per yr), peak	- 0.0087	SE 0.0030, $p = 0.007$ (capacity-weighted reservoir FE, clustered)
capacity-weighted trend (per yr), trough	0.0166	- SE 0.0017, $p = 5.7\text{e-}10$ (capacity-weighted reservoir FE, clustered)
capacity-weighted trend (per yr), amplitude	0.0079	SE 0.0022, $p = 0.00124$ (capacity-weighted reservoir FE, clustered)
capacity-weighted step at 2010, peak	0.1781	- SE 0.0515, $p = 0.00205$ (capacity-weighted reservoir FE, clustered)
capacity-weighted step at 2010, trough	0.1165	- SE 0.0554, $p = 0.0462$ (capacity-weighted reservoir FE, clustered)
capacity-weighted post-2010 drift (per yr), peak	0.0009	SE 0.0011, $p = 0.43$ (capacity-weighted reservoir FE, clustered)
capacity-weighted post-2010 drift (per yr), trough	0.0127	- SE 0.0028, $p = 0.000143$ (capacity-weighted reservoir FE, clustered)
coverage-stable subset share of fleet capacity	0.9770	22 of 25 reservoirs with ≥ 18 band years; totals below use years with all 22 reporting
fleet total peak storage (hm ³), trend and 2010 step	88.0000	- trend $-88.0 \text{ hm}^3/\text{yr}$ ($p = 0.0216$); 2010 step -1549 hm^3 ($p = 0.106$), post-2010 drift $+204.9 \text{ hm}^3/\text{yr}$ ($p = 0.518$); 22 reservoirs with full coverage, 20 of 21 years
fleet total trough (carryover) storage (hm ³), trend and 2010 step	209.2000	- trend $-209.2 \text{ hm}^3/\text{yr}$ ($p = 7.53\text{e-}05$); 2010 step -1254 hm^3 ($p = 0.384$), post-2010 drift $-71.3 \text{ hm}^3/\text{yr}$ ($p = 0.884$); 22 reservoirs with full coverage, 20 of 21 years

Supplementary Table S45 | Hard-balance weight diagnostics and leave-one-out sensitivity. The aridity hard balance achieves SMD 0.000 at the cost of concentrated control weights (maximum normalized share 13%, top five 44%, Kish effective sample 19 of 244), so its stability is tested two ways. The naive leave-one-out drops a top-weighted control with the remaining weights unchanged; because the dropped unit carried part of the aridity balance, the estimate mechanically drifts back toward the confounded primary value, which is expected and is not evidence of fragility. The design-consistent re-balanced leave-one-out re-fits the entropy balance without each unit (log-aridity SMD stays 0.000): the estimate stays near zero for seven of the eight top-weighted controls, but dropping the single most-weighted basin returns the confounded value with balance restored, so the hard-balanced point estimate depends on one control basin. No variant turns the null into detected buffering, since every value in the span lies inside the primary design's permutation null; the hard balance is accordingly read as a fragile corroboration, with the causal weight on the within-region comparison and the upstream placebo, which use no aridity balance.

Quantity	ValueDetail
hard-balance	0.132top-5 share 0.44, Kish ESS 18.9 of 244 controls (participation ratio 0.077)
control weights:	
maximum	
normalized share	
hard-balanced	- primary-design estimate -0.183; hard-balanced perm p > 0.85
ITT slope, all	0.023(Supplementary Table S17)
controls	
naive	0.182estimates span [-0.182, 0.057]; largest excursion dropping unit 0250;
leave-one-out	removing a balancing unit without re-balancing reintroduces the aridity
(weights fixed)	imbalance, so drift toward the confounded -0.183 is mechanical, not
over top-8 controls	fragility
re-balanced	0.184entropy balance re-fit without each unit: estimates span [-0.184, 0.002],
leave-one-out over	log-aridity SMD stays -0.000 to 0.000, ESS 11.6 to 18.1; largest excursion
top-8 controls	dropping unit 0250

Supplementary Table S46 | Routed-network validation of the upstream/downstream classification. The gauge classification is by SRTM elevation, so “upstream” gauges could in principle sit on parallel, hydrologically disconnected tributaries. Every treated gauge and every dam is snapped to its nearest HydroRIVERS v1.0 reach (all 96 gauges snap, median 264 m; all 26 dams, median 124 m) and the NEXT_DOWN topology traced: an upstream gauge is network-connected if its flow path passes its dam’s reach, a downstream gauge if it lies on the dam’s routed downstream path. HydroRIVERS carries a finite drainage threshold, so a gauge that fails the trace mixes true parallel placement with reaches too small for the network; the connected subset is therefore the clean test rather than the unconnected set proof of disconnection. On strictly network-validated gauges the placebo verdict strengthens: the connected-only upstream slope deepens to -0.36 , and the fully validated same-channel contrast is *positive* (down -0.24 , up -0.36 , $D = +0.12$, $p = 0.18$), the unregulated reaches more attenuated than the regulated ones, the opposite sign of a regulation effect.

Quantity	Value Detail
gauges snapped to HydroRIVERS reaches	96.000 of 96 treated gauges with coordinates; median snap distance 264 m; 26 dams snapped (median 124 m); snaps beyond 5000 m excluded
upstream gauges network-connected through their dam	15.000 of 35 upstream gauges: flow path (NEXT_DOWN trace) passes the dam reach; the remainder are parallel tributaries or unsnapped
downstream gauges on the dam’s routed downstream path	16.000 of 61 downstream gauges
upstream placebo slope, network-connected gauges only	- perm $p = 0.265$; contrast vs full downstream panel $D = 0.360 + 0.196$, $p = 0.116$ (1000 perms); raw upstream -0.20
fully network-validated contrast (both classes connected)	0.123 down -0.238 , up -0.360 , perm $p = 0.182$ (1000 perms)

Supplementary Table S47 | Buffering under a smooth national time trend. Year fixed effects absorb the common temporal component of the megadrought, so the primary slope gap is identified from within-year spatial variation in SPEI. Replacing the year fixed effects with a cubic national time trend (month effects retained) restores the temporal drought signal to the identifying variation; every slope is reproduced and remains permutation-null, so the null is not an artifact of the fixed-effects structure or of limited within-year variation.

Quantity	Value	Detail
smooth-trend buffering, ITT	- 0.188	perm $p = 0.320$ (1000 perms); year-FE primary -0.183
smooth-trend buffering, downstream	- 0.170	perm $p = 0.445$ (1000 perms); year-FE primary -0.165
smooth-trend buffering, upstream placebo	- 0.190	perm $p = 0.487$ (1000 perms); year-FE primary -0.201
smooth-trend placebo contrast (down - up)	0.021	cubic national time trend + month FE retain the temporal drought signal that year FE absorb; regulation would give < 0

Supplementary Table S48 | Elevation-partitioned permutation null for the placebo contrast. In the composition-fixed permutation scheme a pseudo-treated control basin contributes an identical gauge set to both panels, so natural up/down heterogeneity is absent from its null contribution, which could compress the null variance of D and make the test anti-conservative (a bias that can only overstate significance, and therefore cannot manufacture the paper’s null). Re-testing under a variant null in which each pseudo-treated control basin’s gauges are partitioned at its median gauge elevation shows no compression: the aggregate null SD is unchanged (0.093 against 0.098) with $p = 0.73$, and the carryover contrast survives the corrected null ($p = 0.025$ against 0.014).

Quantity	Value	Detail
aggregate contrast D under elevation-partitioned null	0.037	perm $p = 0.728$, null SD 0.093 (1000 perms); composition-fixed scheme: $p = 0.718$, null SD 0.098 (Supplementary Table S25)
carryover contrast D under elevation-partitioned null	- 0.232	perm $p = 0.025$, null SD 0.103 (1000 perms); composition-fixed scheme: $p = 0.014$ (Supplementary Table S35)

Supplementary Table S49 | SSI standardized on the 1990–2020 baseline. The primary SSI uses the common 2000–2020 window so every gauge shares one reference period; that window is megadrought-dominated and shorter than the 30-year guideline. Re-standardizing on 1990–2020 (31 years; the CR2 daily record begins in 1900) leaves the estimation window unchanged and extends only the standardization reference: every buffering estimate and the placebo contrast are essentially unchanged, so the transmission slopes are not an artifact of a drought-heavy reference period.

Quantity	Value	Detail
gauges retained under 1990-2020 standardization	171.000	SSI-12 re-standardized per gauge-month (Gringorten) on the 31-year baseline
long-baseline buffering, ITT	- 0.188	perm p = 0.280 (1000 perms); 2000-2020-baseline primary -0.183
long-baseline buffering, downstream	- 0.178	perm p = 0.284 (1000 perms); 2000-2020-baseline primary -0.165
long-baseline buffering, upstream placebo	- 0.198	perm p = 0.458 (1000 perms); 2000-2020-baseline primary -0.201
long-baseline placebo contrast (down - up)	0.020	perm p = 0.893 (1000 perms); regulation would give < 0

Supplementary Table S50 | Drought-conditional streamflow missingness. The no-imputation rule deletes months rather than fills them, so its selection structure is quantified directly. Missingness is genuinely drought-selective (33.8% of drought gauge-months missing against 13.9% otherwise), a truncation of the dry tail shared by treated and control gauges alike, which the differential estimand nets out; what could bias the placebo contrast is a *differential* between gauge classes, and that differential is small and slightly favors upstream loss (32.9% downstream against 35.1% upstream in drought months), the direction that would, if anything, thin the placebo rather than protect it. The dry-tail re-estimates (Supplementary Table S29) and the long-baseline SSI (Supplementary Table S49) carry the same verdict on the retained months.

Quantity	Value	Detail
gauge-month missingness within record span, 2000-2020	0.15670	15.67% of 47124 gauge-months (months failing the ≥ 20 -valid-days rule or absent)
missingness in drought months (SPEI < -1) vs other months	0.33790	drought 33.79% vs non-drought 13.90%
missingness on SPEI (gauge FE, unit-clustered)	- 0.05068	per SPEI unit; SE 0.00891, p = 2.27e-07; a negative sign means drier months are MORE likely missing
drought-month missingness, treated downstream vs upstream gauges	- 0.02160	down 32.92%, up 35.09%; a differential would be needed to bias the contrast D

Supplementary Table S51 | Storage-band trends on the coverage-stable subset. Several reservoirs enter the monitored panel after 2005, and their filling phases are not absorbed by static reservoir fixed effects. Restricting to the 22 reservoirs observed in at least 18 of the 20 band years reproduces the full-fleet pattern: peak and trough decline, no amplitude narrowing, the peak steps down at 2010 with no further drift, and the trough keeps eroding, so mid-panel commissioning does not tilt the trends.

Quantity	Value	Detail
stable-subset trend (per yr), peak	- 0.0134	SE 0.0050, $p = 0.0144$ (22 reservoirs with ≥ 18 band years; full fleet in Supplementary Tables S5, S41)
stable-subset trend (per yr), trough	- 0.0122	SE 0.0027, $p = 0.000177$ (22 reservoirs with ≥ 18 band years; full fleet in Supplementary Tables S5, S41)
stable-subset trend (per yr), amplitude	- 0.0012	SE 0.0057, $p = 0.832$ (22 reservoirs with ≥ 18 band years; full fleet in Supplementary Tables S5, S41)
stable-subset step at 2010, peak	- 0.1935	SE 0.0907, $p = 0.0449$ (22 reservoirs with ≥ 18 band years; full fleet in Supplementary Tables S5, S41)
stable-subset step at 2010, trough	- 0.0743	SE 0.0475, $p = 0.133$ (22 reservoirs with ≥ 18 band years; full fleet in Supplementary Tables S5, S41)
stable-subset post-2010 drift (per yr), peak	- 0.0024	SE 0.0038, $p = 0.532$ (22 reservoirs with ≥ 18 band years; full fleet in Supplementary Tables S5, S41)
stable-subset post-2010 drift (per yr), trough	- 0.0100	SE 0.0037, $p = 0.014$ (22 reservoirs with ≥ 18 band years; full fleet in Supplementary Tables S5, S41)

Supplementary Table S52 | Non-linear aridity-by-forcing sensitivity. Drought transmission need not be linear in aridity, so the residual-aridity adjustment is not left to the linear outcome model: the buffering model adds quadratic aridity-by-forcing interactions. The curvature is strongly detectable (quadratic term $p < 10^{-5}$), and allowing it collapses the apparent intent-to-treat buffering from -0.18 to -0.06 (permutation $p = 0.43$), a third independent route, alongside the within-region estimator and the hard balance, by which flexible aridity adjustment absorbs the apparent effect. This is the signature of a non-linear siting confound, not of regulation.

Quantity	Value	Detail
buffering ITT net of quadratic aridity-by-forcing terms	- 0.057	perm $p = 0.430$ (1000 perms); linear-adjustment primary -0.183
aridity-by-forcing terms (linear; quadratic)	0.033	linear $p = 0.364$; quadratic -0.067, $p = 1.8e-06$ (cluster-robust); the control transmission gradient across aridity is also tercile-checked in Supplementary Table S20

Supplementary Table S53 | Control screening against the national dam inventory.

“Undammed” is defined against the 26 monitored reservoirs, so the matched controls are screened against the full 1,370-impoundment inventory: 37 of 244 matched control basins contain some impoundment, 10 contain a material dam (monitored, registry-large, or wall at least 15 m), and those 10 carry 5.5% of the entropy-balancing weight. Re-fitting with the entropy balance re-estimated on the dam-free control pool leaves every streamflow estimate unchanged, so control-side contamination by unmonitored dams does not attenuate the estimates toward zero.

Quantity	Value Detail
matched control basins containing any inventory dam	37.000 of 244 matched controls; 10 contain a MATERIAL dam (monitored / GRANDE / wall $\geq$ 15 m), holding 5.5% of the entropy-balancing weight
dam-free-control buffering, ITT (re-balanced)	- perm p = 0.353 (1000 perms); primary -0.183; entropy balance 0.183 re-fit on dam-free controls, control ESS 79.6
dam-free-control upstream placebo (re-balanced)	- perm p = 0.472; primary -0.201 0.202
dam-free-control placebo contrast (down - up)	0.037 down -0.165, up -0.202, perm p = 0.700 (1000 perms)

Supplementary Table S54 | Wild-cluster-bootstrap inference for the storage-band models. The storage trends, the 2010 steps, and the post-2010 drifts are estimated with 22 to 26 reservoir clusters, where cluster-robust standard errors can over-reject; a descriptive within-reservoir trend offers no treatment assignment for the design’s randomization inference to permute, so the small-cluster refinement is the restricted (null-imposed) wild cluster bootstrap with Rademacher weights drawn at the reservoir level (2,999 draws per coefficient). Bootstrap p-values are reported beside the analytic ones for every coefficient, unweighted and capacity-weighted. The unweighted results, which carry the main-text claims, tighten under the bootstrap: the peak trend moves from $p = 0.013$ to 0.004 and the 2010 peak step from 0.04 to 0.009, while the amplitude trend and the trough step stay null. The capacity-weighted fits behave oppositely, because weighting concentrates effective weight on a few large reservoirs and so sharpens the small-cluster problem: the weighted trough trend and drift survive ($p = 0.001$) and the weighted amplitude widening survives ($p = 0.031$), but the weighted peak trend ($p = 0.13$) and both weighted 2010 steps ($p = 0.088$ and 0.28) do not, so the analytic p-values reported in Supplementary Table S44 for those three coefficients over-reject and the step evidence rests on the unweighted fit.

Quantity	Value	Detail
unweighted trend (per yr), peak	- 0.0128	analytic cluster-robust $p = 0.0127$; wild cluster bootstrap $p = 0.004$ (1999 draws, Rademacher, null imposed)
unweighted trend (per yr), trough	- 0.0122	analytic cluster-robust $p = 7.33\text{e-}05$; wild cluster bootstrap $p = 0.001$ (1999 draws, Rademacher, null imposed)
unweighted trend (per yr), amplitude	- 0.0006	analytic cluster-robust $p = 0.912$; wild cluster bootstrap $p = 0.984$ (1999 draws, Rademacher, null imposed)
unweighted step at 2010, peak	- 0.1915	analytic cluster-robust $p = 0.04$; wild cluster bootstrap $p = 0.009$ (1999 draws, Rademacher, null imposed)
unweighted step at 2010, trough	- 0.0741	analytic cluster-robust $p = 0.118$; wild cluster bootstrap $p = 0.105$ (1999 draws, Rademacher, null imposed)
unweighted post-2010 drift (per yr), peak	- 0.0024	analytic cluster-robust $p = 0.487$; wild cluster bootstrap $p = 0.513$ (1999 draws, Rademacher, null imposed)
unweighted post-2010 drift (per yr), trough	- 0.0102	analytic cluster-robust $p = 0.00653$; wild cluster bootstrap $p = 0.009$ (1999 draws, Rademacher, null imposed)
capacity-weighted trend (per yr), peak	- 0.0087	analytic cluster-robust $p = 0.007$; wild cluster bootstrap $p = 0.132$ (1999 draws, Rademacher, null imposed)
capacity-weighted trend (per yr), trough	- 0.0166	analytic cluster-robust $p = 5.7\text{e-}10$; wild cluster bootstrap $p = 0.001$ (1999 draws, Rademacher, null imposed)
capacity-weighted trend (per yr), amplitude	- 0.0079	analytic cluster-robust $p = 0.00124$; wild cluster bootstrap $p = 0.031$ (1999 draws, Rademacher, null imposed)
capacity-weighted step at 2010, peak	- 0.1781	analytic cluster-robust $p = 0.00205$; wild cluster bootstrap $p = 0.088$ (1999 draws, Rademacher, null imposed)
capacity-weighted step at 2010, trough	- 0.1165	analytic cluster-robust $p = 0.0462$; wild cluster bootstrap $p = 0.276$ (1999 draws, Rademacher, null imposed)
capacity-weighted post-2010 drift (per yr), peak	- 0.0009	analytic cluster-robust $p = 0.43$; wild cluster bootstrap $p = 0.594$ (1999 draws, Rademacher, null imposed)
capacity-weighted post-2010 drift (per yr), trough	- 0.0127	analytic cluster-robust $p = 0.000143$; wild cluster bootstrap $p = 0.001$ (1999 draws, Rademacher, null imposed)